

C++ Design Patterns: From C++03 to C++17

Fedor G Pikus

Chief Scientist,
Design2Silicon Division

CPPCon 2019


DESIGN PATTERNS

What are Design Patterns?

- [Software] Design pattern is a repeatable, commonly recognized and understood solution to a design problem commonly occurring in software engineering
 - Design problem
 - Commonly occurring
 - Widely accepted solution
 - Known advantages and trade-offs
- Patterns are “design templates”, guidelines for design
- Patterns are compact expressive vocabulary elements
 - Like allusions in speech, they reference commonly known large concepts

What are Design Patterns?

- The “gang of four” (1995) book largely defined the way we think about patterns
- Concept is borrowed from architecture
 - Christopher Alexander, 1977
 - Concept is borrowed from the archetypes
- Software design patterns: Beck and Cunningham, 1987
- Patterns are not static, they evolve
 - Not limited to the “original 23 patterns”
 - Not limited to OOP patterns either
 - Generic programming has patterns too



Should I Always Use Patterns?

- Patterns are very general principles
 - There is always a counter-example where the rigorous application of the rule is worse than breaking it
- A general principle is a "default" rule
 - A guideline that should be followed in absence of a good reason not to
- The majority of everyday work is "not special" and the result is better if this principle is followed
 - The majority of the exceptions (where an idiosyncratic solution is superior) don't gain enough to justify the effort – follow the default rule anyway
- There are new problems never seen before, or new constraints
- There isn't a design pattern for every challenge

Do Design Patterns Depend on the Language?

- Design patterns apply to software design and transcend language
 - In practice, some languages are preferred for certain problems
 - Some languages are more likely to create certain problems
- Some languages offer unique variations on a pattern
 - Strategy pattern: selecting an algorithm for a particular behavior aspect at run time (also known as policy pattern)
 - In C++, more commonly used as policy-based design (compile-time strategy pattern)
 - A pattern could be so hard on a particular language that it's not practical
- Contrast with language-specific idioms
 - Often exist to work around specific problems or deficiencies in a language
 - Change, appear, or disappear as language evolves

Does Language Development Affect Pattern Use?

- Not the same question as “do patterns depend on language?”
- Some patterns are easier to use in a particular language
 - Often, the same overall problem can be solved using multiple designs
- Some patterns just map perfectly to a language feature
 - Null object pattern – `std::optional` (C++17), Maybe (Haskell)
- Language development may change the ease of use balance between patterns
 - Some patterns become “easy/convenient enough” to use widely
- There is friction in the use of the patterns
 - “Small stuff” matters in practice
- A “tipping point” may prompt a different design approach

C++ Evolution and Patterns

- C++ features that significantly reduced the “friction” for using many design patterns:
 - C++14: universal references, variadic templates, lambdas, SFINAE, auto function return types
 - C++17: constructor templates and deduction rules, fold expressions, lambda overloads
- Patterns in this talk are mostly examples to illustrate the effect of C++ evolution on design decision

Hands-On Design Patterns with C++

Solve common C++ problems with modern design patterns and build robust applications



Fedor G. Pikus

Packt
www.packt.com

Restricted © 2019 Mentor Graphics Corporation

Mentor
A Siemens Business

C++ Evolution and Patterns

- C++ features that significantly reduced the “friction” for using many design patterns:
 - C++14: universal references, variadic templates, lambdas, SFINAE, auto function return types
 - C++17: constructor templates and deduction rules, fold expressions, lambda overloads
- Patterns in this talk are mostly examples to illustrate the effect of C++ evolution on design decision



BUILDER

Builder

- The Builder is a pattern used to create (build) objects
- Builder is often used to construct complex objects that are initialized in multiple stages
- Generally, the Builder separates the creation of the object from its representation
 - Builder is a separate class (each class C has its own C_Builder)
 - Separate builder class is also a disadvantage

Builder in C++03

```
class HTMLElement {  
    friend class HTMLBuilder;  
    std::string name_  
    std::string text_  
    std::vector<HTMLElement> children_  
    public:  
    HTMLElement(const std::string& name, const std::string& text)  
        : name_(name), text_(text) {}  
};
```

- Many different ways to write builders

Fluent Builder in C++03

- Many different ways to write builders

```
HTMLElement el = HTMLBuilder("ul").  
    add_child("li", "item 1").  
    add_child("li", "item 2");
```

- Method chaining

Fluent Builder in C++03

- Many different ways to write builders

```
HTMLElement el = HTMLBuilder("ul").
```

```
    add_child("li", "item 1").
```

```
    add_child("li", "item 2");
```

- Method chaining

— Also used to create named arguments to C++ functions:

```
void Fly(double speed, double distance, double angle);
```

```
Fly(5, 180, 2.2);
```

Fluent Builder in C++03

- Many different ways to write builders

```
HTML_Element el = HTMLBuilder("ul").
```

```
    add_child("li", "item 1").
```

```
    add_child("li", "item 2");
```

- Method chaining

— Also used to create named arguments to C++ functions:

```
void Fly(double speed, double distance, double angle);
```

```
Fly(5, 180, 2.2);
```

error-prone
Old boring way

```
Fly(FlyParams().Speed(5).Distance(2.2).Angle(180));
```

New cool way

Fluent Builder in C++03

```
class HTMLBuilder {
    HTMLElement root_;
public:
    explicit HTMLBuilder(const string& name,
        const std::string& text = string()) : root_(name, text) {}
    HTMLBuilder& add_child(const std::string& name,
        const std::string& text) {
        root_.children_.push_back(HTMLElement(name, text));
        return *this;
    }
    operator HTMLElement() const { return root_; }
};
```


Fluent Builder in C++11

```
class HTMLBuilder {  
    HTMLBuilder& add_child(const std::string& name,  
                           const std::string& text) {  
        root_.children_.emplace_back(name, text);  
        return *this;  
    }  
    operator HTMLElement() const { return root_; }  
};
```

■ Usual C++11 enhancements

- Maybe C++11 for-loop to iterate over children

A Very C++11 Builder

```
std::cout << UL{  
    LI{"item 1"},  
    LI{"item 2",  
        UL{LI{"sub-item 2.1"},  
            LI{"sub-item 2.2"}  
        }  
    }  
};
```

- No separate builder class

A Very C++11 Builder

```
class HTMLElement {
    std::string name_;
    std::string text_;
    std::vector<HTMLElement> children_;
public:
    HTMLElement(const std::string& name, const std::string& text)
        : name_(name), text_(text) {}
    HTMLElement(const std::string& name, const std::string& text,
        std::vector<HTMLElement>&& children)
        : name_(name), text_(text), children_(std::move(children)) {}
}; // No more friends (in fact, no HTMLBuilder class)
```

A Very C++11 Builder

- Concrete classes define grammar
- Builder is mostly auto-generated

```
struct UL : public HTMLElement {  
    UL() : HTMLElement("ul", "") {}  
    UL(std::initializer_list<HTMLElement> children) :  
        HTMLElement("ul", "", children) {}  
};
```

A Very C++11 Builder

```
struct LI : public HTML_Element {  
    explicit LI(const std::string& text) : HTML_Element("li", text) {}  
    LI(const string& text, std::initializer_list<HTML_Element> children)  
        : HTML_Element("li", text, children) {}  
};
```

A Very C++11 Builder

```
struct LI : public HTML_Element {
    explicit LI(const std::string& text) : HTML_Element("li", text) {}

    /* Don't I wish...
    LI(const string& text, std::initializer_list<HTML_Element> children)
       : HTML_Element("li", text, children) {} */
    template <typename ... Children>
    LI(const std::string& text, const Children& ... children)
       : HTML_Element("li", text,
                      std::initializer_list<HTML_Element>{children ... }) {}
};
```

Does Language Development Affect Pattern Use?

- C++03 code wasn't particularly bad or annoying...
 - C++11 allows for some improvements (ease of use and efficiency)
- C++11 parameter packs can be used as code generators
 - New twist on the old pattern

VISITOR

Visitor

- The Visitor is a pattern that separates the algorithm from the object structure which is the data for this algorithm.
- Visitor adds new operations to the class hierarchy without modifying the classes themselves
- Open/Closed principle of the software design: a class should be closed for modifications but open for extensions
 - Interface should remain stable under maintenance
 - New functionality can be added to satisfy new requirements
- Useful for public APIs that must be extended by the clients

Visitor – Technical Viewpoint

- Visitor is double dispatch
- Single dispatch:

```
class B {  
    virtual void f() = 0;  
};
```

```
class D1 : public B {  
    void f() { ... do D1 stuff ... };  
};
```

```
class D2 : public B {  
    void f() { ... do D2 stuff ... };  
};
```

```
B* b = ... D1 or D2 ...;
```

```
b->f();
```



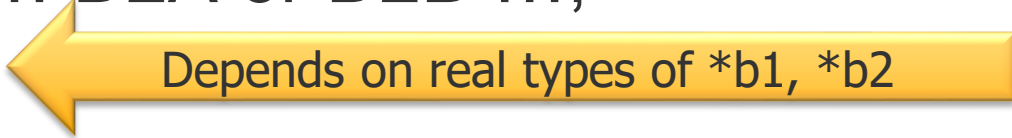
Depends on real type of *b

- The two viewpoints describe the same pattern

Visitor – Technical Viewpoint

- Visitor is double dispatch
- Double dispatch:

```
class B1 {                                class B2 { ... }
    virtual void f(B2*) = 0;
};
class D1A : public B1 { ... }; class D1B : public B1 { ... };
class D2A : public B2 { ... }; class D2B : public B2 { ... };
B1* b1 = ... D1A or D1B ...;
B2* b2 = ... D2A or D2B ...;
b1->f(b2);
```



Depends on real types of *b1, *b2

Why Visitor?

- Public APIs or other cases when changing source is not possible
- A way to keep decision-making decentralized
- Example: serialization
 - Each class knows how to serialize itself
 - There is serialization to disk, buffer, socket, other destination
 - One option is a huge central function with a case for every combination of class and destination (not the only option)
 - Visitor alternative: double dispatch based on class and destination

Classic Visitor (C++03)

- Class hierarchy:

```
class Cat; class Dog;
```

- New operations:

```
class FeedingVisitor {  
    void visit(Cat* c); void visit(Dog* d);  
};
```

- Client code:

```
Cat c("orange"); FeedingVisitor fv;
```

```
c.accept(fv);
```



Double dispatch

Classic C++ Visitor (C++03)

- Class hierarchy:

```
class Pet {  
    std::string color_  
public:  
    Pet(const std::string& color) : color_(color) {}  
    const std::string& color() const { return color_; }  
    virtual void accept(PetVisitor& v) = 0;  
};  
class Cat : public Pet { void accept(PetVisitor& v) { v.visit(this); } };  
class Dog : public Pet { ... };  
Pet* p; p->accept(pv);
```

Depends on the Pet p and the Visitor v

Classic C++ Visitor (C++03)

- Visitors (new operations):

```
class PetVisitor {
    virtual void visit(Cat* c) = 0;
    virtual void visit(Dog* d) = 0;
};

class FeedingVisitor : public PetVisitor {
    void visit(Cat* c) override {
        cout << "Feed tuna to the " << c->color() << " cat" << endl; }
    void visit(Dog* d) override {
        cout << "Feed steak to the " << d->color() << " dog" << endl; }
};
```

Classic C++ Visitor (C++03)

- Client code:

```
FeedingVisitor fv;
```

```
PlayingVisitor pv;
```

```
WalkingVisitor wv;
```

```
Pet* c = new Cat("orange");
```

```
Pet* d = new Dog("brown");
```

```
c->accept(pv);
```

```
d->accept(wv);
```


Why Visitor? And Why Not?

- +** New operations can be added without modifying the hierarchy
 - After the classes are made visitable, once
- +** Impossible to forget to implement an option
 - If the implementation for a class and a visitor type is missing, the code will not compile (pure virtual not overridden)
- Once a class is added, all visitors must be updated
 - Visitor is recommended for “stable hierarchies”
- Visitor does not have privileged access, sacrifices encapsulation
- Visitor functions can take additional arguments and return values, but must be the same types for all visitors
 - Arguments are usually passed to visitors directly

Visitor in Modern C++

- Mostly cleaner and easier to maintain
- Hierarchy has boilerplate visitation code:

```
class Cat : public Pet {  
    void accept(PetVisitor& v) { v.visit(this); }    // Cannot move to Pet  
};
```

- Visitor classes must be declared, have some boilerplate:

```
class FeedingVisitor : public PetVisitor {  
    void visit(Cat* c) override {  
        cout << "Feed tuna to the " << c->color() << " cat" << endl; }  
};
```

Visitor in Modern C++

- Class hierarchy:

```
class Pet { ... }; // Same as before
```

```
template <typename Derived> class Visitable : public Pet {  
    using Pet::Pet;  
    void accept(PetVisitor& v) { v.visit(static_cast<Derived*>(this)); }  
};
```

```
class Cat : public Visitable<Cat> { // Pet is still the base class  
    using Visitable<Cat>::Visitable;  
    ... class-specific code ...  
};
```

- Almost CRTP but not quite

Write once per hierarchy

Boilerplate generator

Visitor in Modern C++

- Visitor and client code:

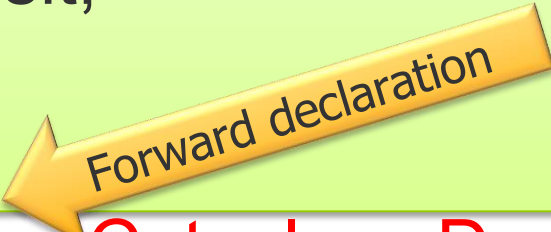
```
auto v(lambda_visitor<PetVisitor>(
    [](Cat* c) { cout << "Let the " << c->color() << " cat out" << endl; },
    [](Dog* d) { cout << "Walk the " << d->color() << " dog" << endl; }));
Pet* p = ...;
p->accept(v);
```

- There is the small matter of implementation
- lambda_visitor<> is written only once (ever)
- PetVisitor is per hierarchy but is auto-generated

Visitor in Modern C++ - Implementation

- PetVisitor is the base class for all visitors in the hierarchy
 - It's essentially a typelist of all visited classes
 - It needs to be updated when a class is added, but in one place only!

```
template <typename ... Types> class Visitor;           // List of classes
template <typename T> class Visitor<T> { virtual void visit(T* t) = 0; };
template <typename T, typename ... Types>
class Visitor<T, Types ...> : public Visitor<Types ...> { // Recursive
    using Visitor<Types ...>::visit;
    virtual void visit(T* t) = 0;
};
using PetVisitor = Visitor<class Cat, class Dog>;
```

 Forward declaration

Write once (ever)

Visitor in Modern C++ - Implementation

- Visitor relies on overload resolution

```
class FeedingVisitor : public PetVisitor {  
    void visit(Cat* c) override;  
    void visit(Dog* d) override;  
};
```

- Lambda visitor uses lambda expressions instead of functions
- Lambda visitor relies on overload resolution of lambda expressions

Visitor in Modern C++ - Implementation

- Visitor relies on overload resolution

```
class FeedingVisitor : public PetVisitor {  
    void visit(Cat* c) override;  
    void visit(Dog* d) override;  
};
```

- Lambda visitor uses lambda expressions instead of functions
- Lambda visitor relies on overload resolution of lambda expressions
- There is no overload resolution of lambda expressions

Divertimento – lambda overload resolution

- The idea is to create a class with overloaded operator():

```
template <typename ... F> struct overload_set : public F ... {  
    overload_set(F&& ... f) : F(std::forward<F>(f)) ... {}  
    using F::operator() ...; // C++17  
};
```

- Also needs a helper function:

```
template <typename ... F> auto overload(F&& ... f) {  
    return overload_set<F ...>(std::forward<F>(f) ...);  
}
```


Divertimento – lambda overload resolution

- Use of the overload set:

```
auto l = overload(  
    [](int i) { std::cout << "i=" << i << std::endl; },  
    [](double d) { std::cout << "d=" << d << std::endl; }  
);  
l(5);  
l(double(5));  
l(float(5));
```

- Can be done in C++14 but does not handle ambiguous overloads as well

Lambda overload resolution in style

- C++17 all the way:

```
template <typename ... F> struct overload_set : public F ... {  
    using F::operator() ...;  
};
```

```
template <typename ... F> overload_set(F&& ... f) ->  
    overload_set<F ...>;
```

- That's not a helper function!
 - No function body, no return
- It's a deduction guide
 - Creates fictional constructors and template deduction rules for them

Lambda overload resolution in style

- Use of the [fancy] overload set:

```
auto l = overload_set{
```

```
    [](int i) { std::cout << "i=" << i << std::endl; },
```

```
    [](double d) { std::cout << "d=" << d << std::endl; },
```



```
};
```

```
l(5);
```

```
l(double(5));
```

```
l(float(5));
```

- Exactly the same as before as far as the client code is concerned

Visitor in Modern C++ - Implementation

- Visitor relies on overload resolution

```
class FeedingVisitor : public PetVisitor {  
    void visit(Cat* c) override;  
    void visit(Dog* d) override;  
};
```

- Lambda visitor uses lambda expressions instead of functions
- Lambda visitor relies on overload resolution of lambda expressions
 - We have it!

Visitor in Modern C++ - Implementation

- We need to iterate over the typelist hidden in PetVisitor
 - We can construct the lambda overload set along the way

```
template <typename Base, typename... > class LambdaVisitor;
```

- Primary template definition, never used

```
template <typename Base, typename T1, typename ... T,  
          typename F1, typename ... F>  
class LambdaVisitor<Base, Visitor<T1, T ...>, F1, F ...>;
```

- Specialization, uses two parameter packs: visitable types T and lambda expressions F

Visitor in Modern C++ - Implementation

- We need to iterate over the typelist hidden in PetVisitor
 - We can construct the lambda overload set along the way

```
template <class Base, class... > class LambdaVisitor;
template <...> class LambdaVisitor<...> :
    private F1, public LambdaVisitor<Base, Visitor<T ...>, F ...> {
    LambdaVisitor(F1&& f1, F&& ... f)
        : F1(std::move(f1)),
          LambdaVisitor<Base, Visitor<T ...>, F ...>(std::forward<F>(f) ...) {}
    void visit(T1* t) override { return F1::operator()(t); }
};
```

- Recursion must end somewhere

Visitor in Modern C++ - Implementation

- Recursion termination – last type in the list:

```
template <typename Base, typename T1, typename F1>
class LambdaVisitor<Base, Visitor<T1>, F1> : private F1, public Base
{
    LambdaVisitor(F1&& f1) : F1(std::move(f1)) {}
    void visit(T1* t) override { return F1::operator()(t); }
};
```

- Only the last class in the inheritance chain inherits from base
 - Base is the PetVisitor class, contains the list of visitable types

Visitor in Modern C++

From the earlier slide

```
using PetVisitor = Visitor<class Cat, class Dog>;
```

```
class Pet { ... }; // Same as before
```

```
template <typename Derived> class Visitable : public Pet {
```

```
    using Pet::Pet;
```

```
    void accept(PetVisitor& v) { v.visit(static_cast<Derived*>(this)); }
```

```
};
```

Write once per hierarchy

```
class Cat : public Visitable<Cat> { using Visitable<Cat>::Visitable; };
```

```
auto v(lambda_visitor<PetVisitor>(
```

```
    [](Cat* c) { cout << "Let the " << c->color() << " cat out" << endl; },
```

```
    [](Dog* d) { cout << "Walk the " << d->color() << " dog" << endl; }));
```

```
Pet* p = ...; p->accept(v);
```


A Very C++17 Visitor

- C++17 has `std::variant` – alternative to polymorphism
- C++17 has `std::visit` – seems like a bold hint

A Very C++17 Visitor

- C++17 has `std::variant` and `std::visit` – we can build a visitor

using Pets = `std::variant`<class Cat, class Dog>;

```
template <typename Visitor, typename Pet>
```

```
void do_visit(const Visitor& v, const Pet& p) {
```

```
    std::visit(v, Pets{p});
```

```
}
```

```
class Cat {
```

```
    Cat(const std::string& color) : color_(color) {}
```

```
    const std::string& color() const { return color_; }
```

```
};
```

```
class Dog { ... also has color() ... };
```


Forward declarations

No visitation interface

Common base not required

A Very C++17 Visitor

- C++17 visitor – the client code

```
auto pv = overloaded { 
    [](const Cat& c) { std::cout << "Drive " << c.color() << " cat nuts"
        " with the laser pointer" << std::endl; },
    [](const Dog& d) { std::cout << "Play fetch with the " << d.color() <<
        " dog" << std::endl; },
};
Cat c("orange");
Dog d("brown");
do_visit(pv, c); // std::visit(pv, Pets{c});
do_visit(pv, d);
```

A Very C++17 Visitor

- C++17 visitor – the client code

```
template <typename Pet> void walk(const Pet& p) {  
    auto v = overloaded {  
        [](const Cat& c) { std::cout << "Let the " << c.color() << " cat out"  
            << std::endl; },  
        [](const Dog& d) { std::cout << "Take the " << d.color() <<  
            " dog for a walk" << std::endl; },  
    };  
    std::visit(v, Pets{p});  
}
```

```
Cat c("orange"); Dog d("brown");  
walk(c); walk(d);
```

A Very C++17 Visitor

- `std::variant` is used instead of object polymorphism
 - No need for a single hierarchy
- Visitation is not routed through the `accept()` method

A Very C++17 Visitor

- `std::variant` is used instead of object polymorphism
 - No need for a single hierarchy
- Visitation is not routed through the `accept()` method
- Harder to write composable visitors:

```
class Family {  
    Cat cat_; Dog dog_;  
    void accept(PetVisitor& v) { cat_.accept(v); dog_.accept(v); }  
};
```

- Serialization/deserialization visitors often use this pattern

Does Language Development Affect Pattern Use?

- C++03 supports the classic OOP visitor (also acyclic visitor)
 - Fair amount of copy-paste boilerplate
 - C++11 removes most of the boilerplate
- C++14 allows visiting lambda expressions
 - With some limitations on overloading (less important for Visitor)
 - Made slightly more compact in C++17
- Definitely much less friction in newer C++ versions
 - Nothing truly radical, but you have to compare with your alternatives
 - Visitor may become easier than the alternative
- C++17 allows visiting variants instead of class hierarchies
 - Has advantages and tradeoffs

SCOPEGUARD

Exception Handling

```
class Record { ... };  
class Database {  
    void insert(const Record& r);  
};
```

- To the caller, `insert()` appears to be a transaction
- It is reasonable to expect transactional behavior
 - `insert()` either succeeds and inserts the record, or fails and nothing happens to the database (exception is thrown)

~~Error Exception Handling~~

```
class Record { ... };  
class Database {  
    int insert(const Record& r);  
};
```

- To the caller, insert() appears to be a transaction
- It is reasonable to expect transactional behavior
 - insert() either succeeds and inserts the record, or fails and nothing happens to the database (exception is thrown)
- It's not about exception handling but error handling


Transactions Are Hard

```
class Database {  
    class Storage { ... }; // Disk storage  
    Storage S;  
    class Index { ... }; // Memory index  
    Index I;  
    void insert(const Record& r) {  
        S.insert(r);  
        I.insert(r);  
    };  
};
```

- The implementation does not guarantee the atomic transaction

Transactions Are Hard

```
class Database {  
    class Storage { ... }; // Disk storage  
    Storage S;  
    class Index { ... }; // Memory index  
    Index I;  
    void insert(const Record& r) {  
        S.insert(r);  
        I.insert(r);  
    };  
};
```



- Nothing happens if the first step fails – so far so good

Transactions Are Hard

```
class Database {  
    class Storage { ... }; // Disk storage  
    Storage S;  
    class Index { ... }; // Memory index  
    Index I;  
    void insert(const Record& r) {  
        S.insert(r);  
        I.insert(r);  
    };
```


Already done!

Fails (throws exception)

- Storage is altered if the second step fails – database corrupted

Transactions Are Easy

```
void Database::insert(const Record& r) {  
    S.insert(r);  
    try { I.insert(r);}  
    catch (...) {  
        S.undo();  
        throw; // Rethrow  
    }  
};
```



- Either all necessary changes are done, or none of them
 - The invariant of the database is maintained
- Exceptions are not required, error codes are handled the same way

Transactions Are Hard

- Undo is hard to do, but easier if you have a point of no return
 - insert() or undo() must be followed by finalize()

```
void Database::insert(const Record& r) {
```

```
    S.insert(r);
```

```
    try { I.insert(r); }
```

```
    catch (...) {
```

```
        S.undo();
```

```
        S.finalize();
```

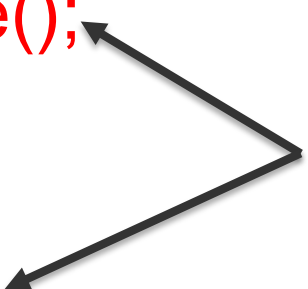
```
        throw;
```

```
    }
```

```
    S.finalize();
```

```
}
```

It's going to be worse if
there are more steps



A Three-Step Transaction

```
if (action1() == SUCCESS) {  
    if (action2() == SUCCESS) {  
        if (action3() == FAIL) {  
            rollback2();  
            rollback1();  
        }  
        cleanup2();  
    } else {  
        rollback1();  
    }  
    cleanup1();  
}
```

- Ugly
- Requires copy-paste
- Gets worse with more steps
- Not composable:
 - The solution for N steps is not the solution for N-1 steps with some more code added
 - To add a step, inner code has to be modified

There Is a Pattern For That

- Resource Acquisition is Initialization (RAII)

```
class StorageFinalizer {  
    StorageFinalizer(Storage& S) : S_(S) {}  
    ~StorageFinalizer() { S_.finalize(); }  
    Storage& S_;  
};
```

- In this case, more like Resource Release is Destruction
 - finalize() happens whenever the finalizer is destroyed

There Is a Pattern For That

- Resource Acquisition is Initialization (RAII)

```
class StorageFinalizer { ... };  
void Database::insert(const Record& r) {  
    S.insert(r);  
    StorageFinalizer SF(S);  
    try { I.insert(r); }  
    catch (...) {  
        S.undo();  
        throw;  
    }  
}
```

- Better, but only so much
 - finalize() is hidden and automated
 - undo() is not

A Three-Step Transaction

```
action1();
Cleanup1 c1; // RAII for cleanup1()
try { action2();
    Cleanup2 c2;
    try { action3(); }
    catch (...) {
        rollback2();
        throw;
    } catch (...) {
        rollback1();
    }
}
```

- Still not composable:
 - rollbackN() is inserted into the innermost scope

A Three-Step Transaction

```
action1();
Cleanup1 c1;
try { action2();
    Cleanup2 c2;
    try { action3(); }
    catch (...) {
        rollback2();
        throw;
    } catch (...) {
        rollback1();
    }
}
```

- Still not composable:
 - rollbackN() is inserted into the innermost scope
- But consider the success (cleanup) path by itself

We're Onto Something here

- Cleanup path is perfectly composable:

```
action1();
```

```
Cleanup1 c1;
```

```
action2();
```

```
Cleanup2 c2;
```

```
...
```

```
actionN();
```

```
CleanupN cN;
```

- Cleanup path is hidden in RAII objects
 - Rollback path is explicit

There Is a Pattern For That

```
class StorageGuard {  
    StorageGuard(Storage& S) : S_(S), commit_(false) {}  
    ~StorageGuard() { if (!commit_) S_.undo(); }  
    void commit() noexcept { commit_ = true; }  
    Storage& S_;  
    bool commit_;  
};
```

- StorageGuard is similar to StorageFinalizer
 - Except the destructor action is conditional
 - Cleanup always happens, rollback happens only on failure

There Is a Pattern For That

```
void Database::insert(const Record& r) {  
    S.insert(r);  
    StorageFinalizer SF(S);    // Arm cleanup action  
    StorageGuard SG(S);      // Arm rollback action (hope to fail?)  
    I.insert(r);  
    SG.commit();              // Disarm rollback if we didn't fail  
}
```

- No try-catch blocks!
- Catch exceptions only to prevent them from propagating
 - Never to execute exception-only code
- Declarative programming (state your intent, and magic happens)

Problems with RAII

```
class StorageGuard {  
    StorageGuard(Storage& S) : S_(S), commit_(false) {}  
    ~StorageGuard() { if (!commit_) S_.undo(); }  
    void commit() noexcept { commit_ = true; }  
    Storage& S_;  
    bool commit_;  
};
```

- RAII classes have to be written for every task
- RAII classes have boilerplate code to capture external variables
- RAII classes are not trivial to write correctly (this one is wrong)

Problems with RAII

```
class StorageGuard {  
    StorageGuard(Storage& S) : S_(S), commit_(false) {}  
    ~StorageGuard() { if (!commit_) S_.undo(); }  
    void commit() noexcept { commit_ = true; }  
    Storage& S_;  
    bool commit_;  
    StorageGuard(const StorageGuard&) = delete;  
    StorageGuard& operator=(const StorageGuard&) = delete;  
};
```

- RAII classes are not trivial to write correctly
 - Really bad things happen if RAII classes are copied

The ScopeGuard

- The ScopeGuard pattern has two elements:
 - The optimal approach to the cleanup/rollback problem – we have already seen it (applies to any deferred action)
 - The recommended implementation – we are about to see it
- This is what we want:

```
S.insert(r);
```

```
ScopeGuard SF(finalize, S);
```

```
ScopeGuard SG(undo, S); // Or something like this
```

```
l.insert(r);
```

```
SG.commit();
```

The ScopeGuard

- The ScopeGuard pattern has two elements:
 - The optimal approach to the cleanup/rollback problem – we have already seen it (applies to any deferred action)
 - The recommended implementation – we are about to see it
- This is what we really want:

```
S.insert(r);
```

```
ScopeGuardFail SG(undo, S); // Or something like this
```

```
ScopeGuardSuccess SF(finalize, S);
```

```
I.insert(r);
```

```
// No explicit SG.commit() – committed if no exceptions thrown
```

ScopeGuard in C++03 – Client code

```
S.insert(r);  
const ScopeGuardImplBase& SG = MakeGuard(undo, S);  
const ScopeGuardImplBase& SF = MakeGuard(finalize, S);  
I.insert(r);  
SG.commit();
```

- Explicit commit
- MakeGuard helper function
 - Fixed number of arguments (one object, one function)
- Cheating alert: this implementation calls a non-member function

```
void undo(Storage& S) { S.undo(); }
```

ScopeGuard in C++03 - Implementation

```
template <typename Func, typename Arg>
ScopeGuardImpl<Func, Arg> MakeGuard(const Func& f, Arg& arg) {
    return ScopeGuardImpl<Func, Arg>(f, arg);
}
```

- Cheating alert: this implementation calls a non-member function

```
void undo(Storage& S) { S.undo(); }
```

- Member function guard is possible but the syntax is even more verbose

ScopeGuard in C++03 - Implementation

```
template <typename Func, typename Arg>
class ScopeGuardImpl : public ScopeGuardImplBase {
    public:
        ScopeGuardImpl(const Func& f, Arg& arg) : func_(f), arg_(arg) {}
        ~ScopeGuardImpl() { if (!commit_) func_(arg_); }
    private:
        const Func& func_;
        Arg& arg_;
};
```

ScopeGuard in C++03 - Implementation

```
class ScopeGuardImplBase {
public:
    ScopeGuardImplBase() : commit_(false) {}
    void commit() const throw() { commit_ = true; }
protected:
    ScopeGuardImplBase(const ScopeGuardImplBase& other)
        : commit_(other.commit_) { other.commit(); }
    ~ScopeGuardImplBase() {}
    mutable bool commit_;
private:
    ScopeGuardImplBase& operator=(const ScopeGuardImplBase&);
};
```

ScopeGuard in C++1X

- C++11: real move constructor, variadic templates
- C++14: return type deduction in functions
- C++17: template type deduction in constructors
- It's still ugly and somewhat limiting, there is a better way:
- Lambda expressions!
 - ScopeGuard pattern is radically changed by C++11 and again C++17

ScopeGuard in C++11/14

- Instead of writing RAII classes we can use lambda expressions:

```
S.insert(r);
```

```
auto SF = MakeGuard([&] { S.finalize(); });
```

```
auto SG = MakeGuard([&] () { S.undo(); }); // () is optional
```

```
I.insert(r);
```

```
SG.commit(); // Still explicit commit
```

- Automatic variable capture (S)
- Any cleanup/rollback code, no limitations on the number of arguments or types of functions to call

ScopeGuard in C++17

- Instead of writing RAII classes we can use lambda expressions:

```
S.insert(r)
```

```
ScopeGuard SF([&] { S.finalize(); });
```

```
ScopeGuard SG([&] { S.undo(); });
```

```
I.insert(r);
```

```
SG.commit(); // Optional
```

- No MakeGuard function
- Automatic success/failure detection is possible
 - Only if failure means exception and success means return (any return)

ScopeGuard in C++1x - Implementation

```
class ScopeGuardBase {  
    public:  
    ScopeGuardBase() : commit_(false) {}  
    void commit() noexcept { commit_ = true; } // Not const now  
    protected:  
    ScopeGuardBase(ScopeGuardBase&& other) // Real move ctor!  
        : commit_(other.commit_) { other.commit(); }  
    ~ScopeGuardBase() {}  
    bool commit_; // Not mutable anymore  
    ScopeGuardBase& operator=(const ScopeGuardBase&) = delete;  
};
```

ScopeGuard in C++1x - Implementation

```
template <class Func> class ScopeGuard : public ScopeGuardBase {
public:
    ScopeGuard(Func&& func) : func_(func) {}
    ScopeGuard(const Func& func) : func_(func) {}
    ~ScopeGuard() { if (!commit_) func_(); }
    ScopeGuard(ScopeGuard&& other)
        : ScopeGuardBase(std::move(other)), func_(other.func_) {}
private:
    Func func_;
};
```

ScopeGuard in C++11/14 - Implementation

```
template <typename Func>
ScopeGuard<Func> MakeGuard(Func&& func) {
    return ScopeGuard<Func>(std::forward<Func>(func));
}
```

- In C++17 the factory function is not needed
 - Constructors deduce template parameters

ScopeGuard in C++1X

```
action1();  
ScopeGuard cleanup1([&] { ... });  
ScopeGuard rollback1([&] { ... });  
action2();  
ScopeGuard cleanup2([&] { ... });  
ScopeGuard rollback2([&] { ... });  
action3();  
rollback1.commit();  
rollback2.commit();
```

- Composable ScopeGuard – just add actionN(), commit/rollback guards, and commit() at the end

ScopeGuard and Exceptions

- In C++, only one exception can propagate at any time
 - Not “only one exception can be throw at any time”

```
void exception_handler(...) { // Called if exception X is thrown
    try { ... throw 1; } catch ( int ) { };
} // Only exception X is propagating - OK
```

- This presents problems for ScopeGuard:
 - Rollback (undo) must not throw (if action fails and undo fails, then what?)
 - If rollback throws, we either die, or ignore the exception (shielded guard)
- Commit also must not throw, otherwise false rollback will run
 - This part is easy, it's just a flag set

ScopeGuard and Exceptions

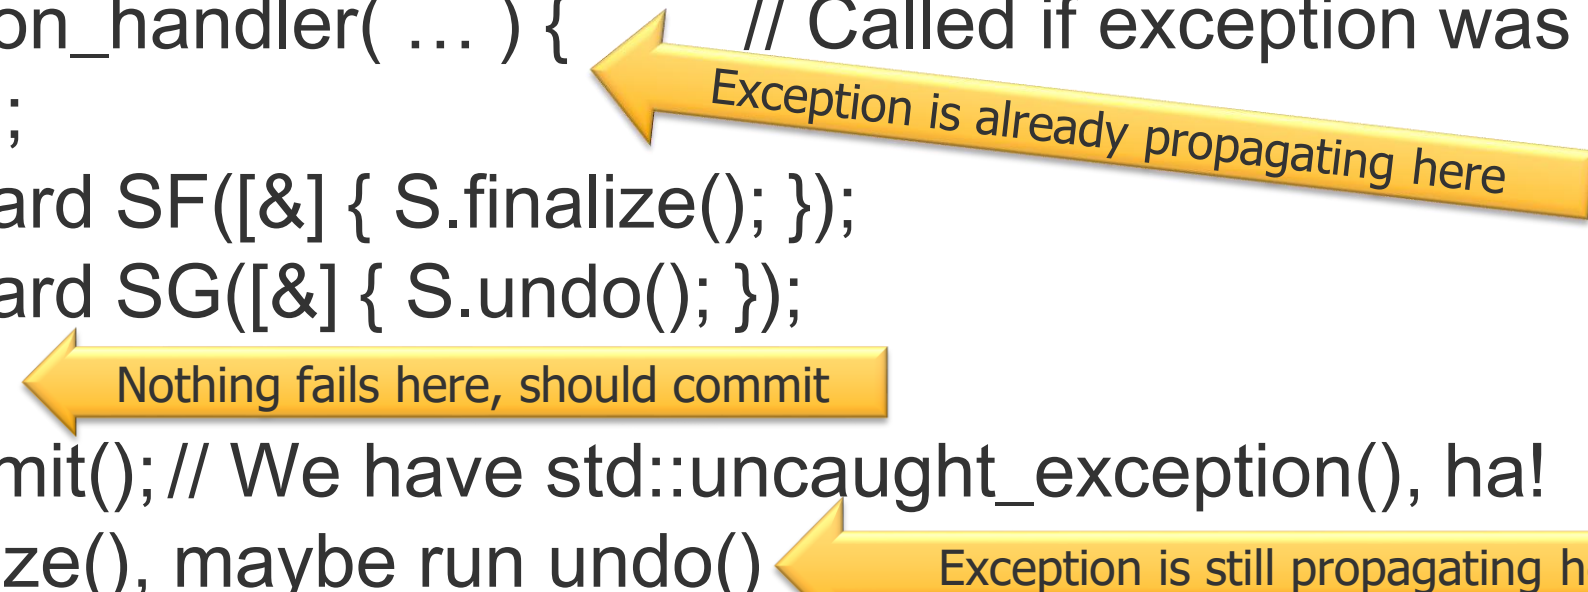
```
{  
  S.insert(r);  
  ScopeGuard SF([&] { S.finalize(); });  
  ScopeGuard SG([&] { S.undo(); });  
  I.insert(r);  
  SG.commit(); // Why?!  
} // run finalize(), maybe run undo()
```



- `std::uncaught_exception()` – return true iff exception is currently propagating
 - Almost enough to auto-detect commit

ScopeGuard and Exceptions

```
void exception_handler( ... ) { // Called if exception was thrown
    S.insert(r);
    ScopeGuard SF([&] { S.finalize(); });
    ScopeGuard SG([&] { S.undo(); });
    I.insert(r);
    //SG.commit(); // We have std::uncaught_exception(), ha!
} // run finalize(), maybe run undo()
```



- `std::uncaught_exception()` – return true iff exception is currently propagating
 - If exception was already propagating when a ScopeGuard was armed, rollback will occur even when guarded actions succeeded

ScopeGuard and Exceptions in C++17

```
void exception_handler( ... ) { // Called if exception was thrown
    S.insert(r);                // N exceptions are already propagating here
    ScopeGuard SF([&] { S.finalize(); });
    ScopeGuardFail SG([&] { S.undo(); });
    I.insert(r);                // Nothing fails here, should commit
} // run finalize(), maybe run undo() // Still N exceptions – success!
```

- `std::uncaught_exceptions()` – return the count of exceptions currently propagating
 - If new exception was thrown since ScopeGuard was armed, it's a failure

ScopeGuard and Exceptions in C++17

```
class UncaughtExceptionDetector {  
    const int c_;  
public:  
    UncaughtExceptionDetector() : c_(std::uncaught_exceptions()) {}  
    operator bool() const noexcept {  
        return std::uncaught_exceptions() > c_;  
    }  
};
```

- Helper class to count exceptions

ScopeGuard and Exceptions in C++17

```
template <typename Func> class ScopeGuardFail {  
    Func func_;  
    UncaughtExceptionDetector detector_;  
public:  
    ScopeGuardFail(Func&& func) : func_(func) {}  
    ScopeGuardFail(const Func& func) : func_(func) {}  
    ~ScopeGuardFail() { if (detector_) func_(); }  
    ScopeGuardFail(ScopeGuardFail&& other) : func_(other.func_) {}  
};
```

- Scope guard for rollback and other failure handling actions

ScopeGuard and Exceptions in C++17

```
template <typename Func> class ScopeGuard {  
    Func func_;  
    public:  
    ScopeGuard(Func&& func) : func_(func) {}  
    ScopeGuard(const Func& func) : func_(func) {}  
    ~ScopeGuard() { func_(); }  
    ScopeGuard(ScopeGuard&& other) : func_(other.func_) {}  
};
```

- Scope guard for cleanup and other actions that always happen

ScopeGuard and Exceptions in C++17

```
{  
    S.insert(r);  
    ScopeGuard SF([&] { S.finalize(); });           // Always happens  
    ScopeGuardFail SG([&] { S.undo(); });         // If I.insert() throws  
    I.insert(r);  
}
```

- Composable – just add action and guards
- Client code is as simple as it gets
 - But only if success == exception
- Would this be enough to make you change your error handling to use exceptions only?

Does Language Development Affect Pattern Use?

- C++03 has just enough clever hacks to implement a mostly working ScopeGuard
 - No major problems but a lot of minor annoyances (i.e. friction)
- C++14 removes most of the friction by using lambda expressions
- C++17 supports automatic detection of success or failure on exit
 - Only if exceptions are used for any failure
- The pattern is much easier to use, may influence design decisions

STRATEGY

Strategy Pattern

- Enables run-time selection of a specific algorithm for a particular behavior
- Also known as the Policy Pattern
- In C++ is mostly used at compile-time

POLICY-BASED DESIGN

Policy-Based Design

```
template <class T, class DeletePolicy, class CopyPolicy,  
        class MovePolicy, class DebugPolicy> class SuperSmartPtr { ... };
```

- Each policy controls the specific behavior aspect
- Each policy can have a default
 - Changing 15th policy requires repeating all preceding defaults

Policy-Based Design

```
template <class T, class DeletePolicy, class CopyPolicy,  
        class MovePolicy, class DebugPolicy> class SuperSmartPtr { ... };
```

- Each policy controls the specific behavior aspect
- Each policy can have a default
 - Changing 15th policy requires repeating all preceding defaults
- Wait, what?! 15th? – policy customization aspects tend to increase
- In practice, everyone needs a small subset of policy options
 - Typical death spiral of policy-based designs: number of policies grows until everything has a policy, then nobody uses policy-based types because of all the policies they don't need but have to specify
- Solution is aliases for policy types

Policy Aliases in C++03

```
template <class T, class DebugPolicy> class MyPtr :  
    public SuperSmartPointer<T, ArrayDeleter, NoCopy,  
                             MoveOK, DebugPolicy> {  
    MyPtr(T* p) ...;  
    MyPtr(const MyPtr&) ... and MyPtr&& ... and more  
};
```

- All constructors must be repeated in all derived classes
 - Seems trivial, in practice may be enough friction to make alternatives preferable

Policy Aliases in C++1X

```
template <class T, class DebugPolicy> class MyPtr :  
    public SuperSmartPtr<T, ArrayDeleter, NoCopy,  
                        MoveOK, DebugPolicy> {  
    Using SuperSmartPtr<... all template args ...>::SuperSmartPtr;  
};
```

- All constructors can be “resurrected” at once
- Template arguments must be repeated

Policy Aliases in C++1X

```
template <class T, class DebugPolicy>  
using MyPtr = SuperSmartPtr<T, ArrayDeleter, NoCopy,  
                          MoveOK, DebugPolicy>;
```

- Nothing to repeat, very compact
- But no C++17 constructor argument deduction for template aliases

```
MyPtr p(new A, VeryVerboseLogger()); // Does not work
```

Does Language Development Affect Pattern Use?

- Policy-based design works in C++03
 - Trade-offs and drawbacks are known but not always avoided
 - In practice, policy-based design tend to evolve toward unmanageable complexity
 - Sometimes simple commonly used classes are reimplemented even when they are really a particular case of a 15-parameter policy template
- C++14/17 enhancements are not major but practically significant
 - Some of the drawbacks become less severe
 - Complexity is easier to manage

EVOLUTION OF LANGUAGE AND OF DESIGN PATTERNS

Do Design Patterns Depend on the Language?

Does Language Development Affect Pattern Use?

- Design patterns apply to software design and transcend language
- Design pattern drawbacks/trade-offs are more language-dependent
- Drawbacks are not to be analyzed in abstract
 - You have a problem, you need a solution, you'll have to choose one
 - Each solution has some drawbacks and trade-offs
- Language development helps some patterns more than others
 - There is always friction in the use of the patterns (price of complexity)
 - “Small stuff” done many times every day matters in practice
- A “tipping point” in the balance of different patterns may lead to a completely different design

Mentor[®]
A Siemens Business

www.mentor.com